

CloudGuard

Cloud security, evidence ingestion, and GRC dashboard

Project deck maintained in GitHub with Marp, exported to HTML, PDF, and PPTX.

Project Objective

Build a containerized security platform that collects scanner output and turns it into compliance-ready evidence.

- Scan AWS, GCP, Kubernetes, and Infrastructure-as-Code targets
- Keep scanner workloads decoupled from the portal
- Normalize findings through evidence connectors
- Store evidence artifacts with checksum, source, control mapping, and metadata
- Show leadership and audit teams one GRC-oriented view

What The Platform Includes


Area	Current project capability
Cloud security	AWS, GCP, CloudFox-assisted cloud posture and attack-path scanning
Kubernetes	Kubeconfig-based live scan, manifest scan, RBAC, services, namespaces, network policies
IaC	Uploaded or repository files for Terraform and CloudFormation
App security	OWASP ZAP, Nuclei, Gitleaks, Trivy, Safety integration paths
GRC	Compliance dashboard, control evidence, source summary, audit artifacts
AI	Optional OpenAI-compatible endpoint for analysis and recommendations

Evidence Architecture

ARCHITECTURE

Evidence pipeline into GRC

Every source sends normalized evidence; the GRC layer maps it to controls, risk and remediation.



Decoupled Scan Job Flow

```
Dashboard / API
  |
  v
POST /api/jobs/scan
  |
  v
PostgreSQL scan_jobs queue
  |
  v
scan-worker on VM
  |
  v
scan results + evidence artifacts
  |
  v
Compliance dashboard + reports
```

The portal queues and displays work. Scanner workers process jobs separately so heavy scans do not block the UI.

Evidence Connector Contract

External tools do not need to be inside the dashboard. They can send normalized evidence to one API.

```
POST /api/evidence
x-connector-token: <CONNECTOR_TOKEN>
x-cloudguard-user: <user-or-tenant>
```

```
{
  "control_id": "A.1.3.7",
  "source_system": "trivy",
  "scanner_type": "container",
  "artifact_type": "json",
  "payload": { "findings": [] },
  "metadata": { "connector": "trivy-file" }
}
```

Data Stored For Evidence

Table	Purpose
scan_jobs	Job ID, user, providers, status, attempts, worker ID, result, error, timestamps
evidence_artifacts	Evidence ID, job ID, user, control ID, source system, scanner type, storage URI, checksum, payload

Operational status supports:

- queued , running , completed , failed , dead_letter , cancelled
- retry/cancel APIs for operations
- latest evidence and control-level evidence lookup

Kubernetes Scanner Coverage

The Kubernetes scanner now supports both live cluster scanning and manifest scanning.

Mode	What it checks
Live kubeconfig	Pods, services, secrets, namespaces, deployments, jobs, cronjobs
RBAC	RoleBindings and ClusterRoleBindings with risky permissions
Network	LoadBalancer/NodePort exposure and missing NetworkPolicies
Managed clusters	EKS, GKE, AKS metadata enrichment when cloud credentials exist
Drift	Compare live resources with Kubernetes IaC manifests

IaC Scanner Coverage

The IaC scanner accepts full files, so users can upload repository snippets or complete templates.

IaC type	Detection examples
Terraform	Public S3 ACL, missing public access block, unrestricted security groups
Terraform	Public databases, missing RDS encryption, public IPs, wildcard IAM
CloudFormation	Public buckets, missing encryption, public ingress, broad IAM
Uploaded files	<code>.tf</code> , <code>.tfvars</code> , <code>.hcl</code> , <code>.yaml</code> , <code>.yml</code> , <code>.json</code>


Uploaded scans also persist evidence with `control_id = iac.uploaded_files` .

Deployment PoC

SETUP

Docker/Dokploy PoC deployment model

The GRC and scanner services can run as separate Docker services while Dokploy handles HTTPS, domains and logs.



PoC VM: 4 vCPU / 8 GB RAM / 100 GB SSD recommended for scanner-heavy demos.

Recommended Deployment Shape

Layer	Recommended placement
Portal and GRC	VM/Dokploy or DigitalOcean app layer
PostgreSQL	Managed DB or private VM database for PoC
Object storage	MinIO for PoC, S3-compatible storage for production
Scanner workers	Separate VM on E2E or other economical provider
Evidence connectors	Separate service/VM posting into <code>/api/evidence</code>

Recommended PoC VM for Docker-heavy single-node validation: **4 vCPU, 8 GB RAM, 100 GB SSD**.

Multi-Tenant Direction

The current implementation separates data primarily by `user_id`.

For production multi-tenant architecture, scanners and connectors should send:

- `tenant_id` for organization-level segregation
- `user_id` for operator/auditor attribution
- `source_system`, `scanner_type`, `control_id`, and `job_id` for traceability
- tenant-scoped credentials, jobs, findings, evidence, reports, and audit events

This lets every scanner follow the same evidence contract without coupling to the portal internals.

What This PoC Proves

- CloudGuard can run as a Docker Compose stack through Dokploy
- Dashboard, API, workers, database, and evidence storage can run together for PoC
- Scanner jobs are separated from the portal through `scan_jobs`
- Evidence connectors can ingest external scanner output
- GRC dashboard can map technical evidence to control status
- The platform can be hosted away from AWS while still scanning AWS accounts

GitHub + Marp Workflow

```
graph TD; A[Edit this Markdown deck] --> B[Commit / PR on GitHub]; B --> C[GitHub Action builds HTML + PDF + PPTX]; C --> D[GitHub Pages shares latest project/workshop deck];
```

Edit this Markdown deck
↓
Commit / PR on GitHub
↓
GitHub Action builds HTML + PDF + PPTX
↓
GitHub Pages shares latest project/workshop deck

Same source file, multiple outputs, easy review history.

Recommendation

Use this Marp setup for the workshop and project updates.

Keep PPTX/PDF exports for formal sharing, and publish the HTML deck through GitHub Pages for live updates.